

Bios 6301: Assignment 2

55/55



AUTHOR

Yiqing Pan

Due Tuesday, 17 September, 1:00 PM

50 points total.

Add your name as `author` to the file's metadata section.

Submit a single quarto file (named `homework2.qmd`) by email to huiding.chen@vanderbilt.edu. Place your R code in between the appropriate chunks for each question. Check your output by using the `Render` button in RStudio.

1. **Working with data** In the `datasets` folder on the course GitHub repo, you will find a file called `cancer.csv`, which is a dataset in comma-separated values (csv) format. This is a large cancer incidence dataset that summarizes the incidence of different cancers for various subgroups. (18 points)

```
library(dplyr)
```

载入程序包: 'dplyr'

The following objects are masked from 'package:stats':

`filter, lag`

The following objects are masked from 'package:base':

`intersect, setdiff, setequal, union`

1. Load the data set into R and make it a data frame called `cancer.df`. (2 points)

```
cancer.df <- read.csv("cancer.csv")
```

2. Determine the number of rows and columns in the data frame. (2)

```
nrow(cancer.df); ncol(cancer.df)
```



```
[1] 42120
```

```
[1] 8
```

3. Extract the names of the columns in `cancer.df`. (2)

```
colnames(cancer.df)
```

```
[1] "year"      "site"      "state"     "sex"       "race"  
[6] "mortality" "incidence" "population"
```

4. Report the value of the 3000th row in column 6. (2)

```
cancer.df[3000, 6]
```

```
[1] 350.69
```

5. Report the contents of the 172nd row. (2)

```
cancer.df[172, ]
```

```

year              site state sex race mortality incidence
172 1999 Brain and Other Nervous System nevada Male Black      0      0
population
172      73172

```

6. Create a new column that is the incidence *rate* (per 100,000) for each row. The incidence rate is the ``(number of cases)/(population at risk)``, which in this case means ``(number of cases)/(population at risk) * 100,000``. (3)

```

cancer.df <- cancer.df %>% mutate(incidence_rate = incidence/population *100000)

head(cancer.df, 10)

```

```

year              site state sex race mortality
1 1999 Brain and Other Nervous System alabama Female Black      0.00
2 1999 Brain and Other Nervous System alabama Female Hispanic 0.00
3 1999 Brain and Other Nervous System alabama Female White    83.67
4 1999 Brain and Other Nervous System alabama Male Black      0.00
5 1999 Brain and Other Nervous System alabama Male Hispanic 0.00
6 1999 Brain and Other Nervous System alabama Male White   103.66
7 1999 Brain and Other Nervous System alaska Female Black      0.00
8 1999 Brain and Other Nervous System alaska Female Hispanic 0.00
9 1999 Brain and Other Nervous System alaska Female White    0.00
10 1999 Brain and Other Nervous System alaska Male Black      0.00
incidence population incidence_rate
1      19      623475      3.047436
2       0       28101      0.000000
3     110     1640665      6.704598
4      18     539198      3.338291
5       0      37082      0.000000
6     145     1570643      9.231888
7       0      12710      0.000000
8       0      11664      0.000000
9       0     220036      0.000000
10      0      13900      0.000000

```

7. How many subgroups (rows) have a zero incidence rate? (2)

```
nrow(cancer.df %>% filter(incidence_rate == 0))
```

```
[1] 23191
```

8. Find the subgroup with the highest incidence rate.(3)

```
cancer.df[which(cancer.df$incidence_rate == summary(cancer.df$incidence_rate)[6]), ]
```

	year	site	state	sex	race	mortality	incidence
5797	1999	Prostate	district of columbia	Male	Black	88.93	420
		population	incidence_rate				
5797		160821	261.1599				

2. Data types (10 points)

1. Create the following vector: `x <- c("5", "12", "7")`. Which of the following commands will produce an error message? For each command, Either explain why they should be errors, or explain the non-erroneous result. (4 points)

```
x <- c("5", "12", "7")
```

```
max(x)
```

```
[1] "7"
```

```
sort(x)
```

```
[1] "12" "5"  "7"
```

```
#sum(x)
```

Command `sum(x)` will produce an error message, indicating that `x` is a invalid 'type' (character).

`max(x)` is a function that returns the maxima of the input values. It works since the characters will be coerced to numeric values and then the maximum will be computed.

`sort(x)` is a function that orders a vector into ascending or descending order. It can work since characters can be easily sorted lexicographically.

`sum(x)` is a function that returns the sum of all the values present in the arguments.`sum(x)` enters an error message since it requires numeric input from the start and does not handle implicit coercion. However, by adding `"` to each numeric value here, we made `x` a character vector.

2. For the next two commands, either explain their results, or why they should produce errors. (3 points)

```
...
y <- c("5", 7, 12)
y[2] + y[3]
...
```

This produces an error since one of the values we are trying to add up is non-numeric. We can not use non-numeric in arithmetic operation.

3. For the next two commands, either explain their results, or why they should produce errors. (3 points)

```
z <- data.frame(z1="5",z2=7,z3=12)
z[1,2] + z[1,3]
```

```
[1] 19
```

```
```
```

```
z <- data.frame(z1="5",z2=7,z3=12)
z[1,2] + z[1,3]
```
```

With these 2 lines, we can get 19 as a result. This works since z is created as a dataframe with 1 row and 3 columns, in which z[1,2] refers to the value stored in the second column of the first row(numeric value 7) and z[1,3] refers to the value stored in the third column of the first row(numeric value 12). We can add up 2 numerical values using +, and the result is simply the addition of 7 and 12, which equals 19.

3. **Data structures** Give R expressions that return the following matrices and vectors (*i.e.* do not construct them manually). (3 points each, 12 total)

1. (1, 2, 3, 4, 5, 6, 7, 8, 7, 6, 5, 4, 3, 2, 1)

```
c((1:8), (7:1))
```

```
[1] 1 2 3 4 5 6 7 8 7 6 5 4 3 2 1
```

2. \$(1,2,2,3,3,3,4,4,4,4,5,5,5,5,5)\$

```
rep(1:5, times = 1:5)
```

```
[1] 1 2 2 3 3 3 4 4 4 4 5 5 5 5 5
```

3. $\begin{pmatrix} 0 & 1 & 1 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{pmatrix}$

```
0 & 1 & 1 \\
1 & 0 & 1 \\
1 & 1 & 0 \\
\end{pmatrix}
```

```
matrix(1, 3, 3) - diag(1, 3, 3)
```

```
      [,1] [,2] [,3]
[1,]    0    1    1
[2,]    1    0    1
[3,]    1    1    0
```

```
4.  $\begin{pmatrix} 1 & 2 & 3 & 4 \\ 1 & 4 & 9 & 16 \\ 1 & 8 & 27 & 64 \\ 1 & 16 & 81 & 256 \\ 1 & 32 & 243 & 1024 \end{pmatrix}$ 
```



```
cbind(
  rep(1, 5),
  2^(1:5),
  3^(1:5),
  4^(1:5)
)
```

```
      [,1] [,2] [,3] [,4]
[1,]    1    2    3    4
[2,]    1    4    9   16
[3,]    1    8   27   64
[4,]    1   16   81  256
[5,]    1   32  243 1024
```

4. Basic programming (10 points)

1. Let $h(x, n) = 1 + x + x^2 + \dots + x^n = \sum_{i=0}^n x^i$. Write an R program to calculate $h(x, n)$ using a `for` loop. As an example, use `x = 5` and `n = 2`. (5 points)

```
x = 5
n = 2
h = 1
for(i in (1:n)) {
  h = h + x^i
}

h
```

```
[1] 31
```

2. If we list all the natural numbers below 10 that are multiples of 3 or 5, we get 3, 5, 6 and 9. The sum of these multiples is 23. Write an R program to perform the following calculations. (5 points)

1. Find the sum of all the multiples of 3 or 5 below 1,000. (3, [euler1] (<https://projecteuler.net/problem=1>))

```
multi_3 <- NULL
multi_5 <- NULL

for (i in (1:999)) {
  if (i%3 == 0) {
    multi_3 <- c(multi_3, i)
  } else if (i%5 == 0) {
```

```
    multi_5 <- c(multi_5, i)
  }
}

sum(multi_3, multi_5)
```

[1] 233168

2. Find the sum of all the multiples of 4 or 7 below 1,000,000. (2)

```
sum_multi_4or7 = 0

for (i in (1:999999)) {
  if ((i%%4 == 0)|(i%%7 == 0)) {
    sum_multi_4or7 = sum_multi_4or7 + i
  }
}

sum_multi_4or7
```

[1] 178571071431

3. Each new term in the Fibonacci sequence is generated by adding the previous two terms. By starting with 1 and 2, the first 10 terms will be \$(1, 2, 3, 5, 8, 13, 21, 34, 55, 89)\$. Write an R program to calculate the sum of the first 15 even-valued terms. (5 bonus points, [euler2] (<https://projecteuler.net/problem=2>))

```
even_val <- c(2)
i <- 1

a = 1
b = 2

while (i < 15) {
  new_a = b
  b = a+b
  a = new_a
  if (b %% 2 == 0) {
    i = i + 1
    even_val[i] = b
  }
}

sum(even_val)
```

[1] 1485607536

Some problems taken or inspired by projecteuler.