

# Bios 6301: Assignment 7

44

Yiqing Pan

*Due Thursday, 31 October, 1:00 PM*

$5^{n=\text{day}}$  points taken off for each day late.

40 points total.

Submit a single quarto file (named `homework7.qmd`), along with a valid PDF output file. Inside the file, clearly indicate which parts of your responses go with which problems (you may use the original homework document as a template). Add your name as `author` to the file's metadata section. Raw R code/output or word processor files are not acceptable.

Failure to name file `homework7.qmd` or include author name may result in 5 points taken off.

```
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'  
  
## The following objects are masked from 'package:stats':  
##  
##   filter, lag  
  
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

## Question 1

**21 points**

Use the following code to generate data for patients with repeated measures of A1C (a test for levels of blood glucose).

```
genData <- function(n) {  
  if(exists(".Random.seed", envir = .GlobalEnv)) {  
    save.seed <- get(".Random.seed", envir = .GlobalEnv)  
    on.exit(assign(".Random.seed", save.seed, envir = .GlobalEnv))  
  } else {  
    on.exit(rm(".Random.seed", envir = .GlobalEnv))  
  }  
  set.seed(n)  
  subj <- ceiling(n / 10)  
  id <- sample(subj, n, replace=TRUE)
```

```

times <- as.integer(difftime(as.POSIXct("2005-01-01"), as.POSIXct("2000-01-01"), units='secs'))
dt <- as.POSIXct(sample(times, n), origin='2000-01-01')
mu <- runif(subj, 4, 10)
a1c <- unsplit(mapply(rnorm, tabulate(id), mu, SIMPLIFY=FALSE), id)
data.frame(id, dt, a1c)
}
x <- genData(500)

```

Perform the following manipulations: (3 points each)

1. Order the data set by id and dt.

```
head(x, 10)
```

```

##      id          dt      a1c
## 1  39 2001-11-30 00:36:10 8.758993
## 2  27 2000-07-20 17:05:06 8.233413
## 3  47 2001-02-06 18:40:44 3.950582
## 4  50 2001-04-05 01:55:44 6.707085
## 5  41 2003-08-05 18:30:14 4.047338
## 6  31 2000-12-16 10:38:26 9.357902
## 7  44 2004-01-09 11:17:40 7.967202
## 8  47 2004-10-12 11:27:02 4.559074
## 9  42 2001-10-02 14:04:39 6.932304
## 10 18 2000-08-23 02:44:30 6.122560

```

```

order_x <- x[order(x$id, x$dt), ]
rownames(order_x) <- 1:nrow(order_x)

head(order_x, 10)

```

```

##      id          dt      a1c
## 1    1 2001-05-08 16:22:52 7.309995
## 2    1 2001-06-17 22:42:23 8.310721
## 3    1 2001-08-17 16:51:46 6.548845
## 4    1 2001-12-14 14:50:29 5.985275
## 5    1 2002-08-19 13:51:47 6.011547
## 6    1 2003-03-22 03:51:36 7.243858
## 7    1 2003-06-27 01:01:34 5.170870
## 8    2 2001-03-05 22:24:43 9.237660
## 9    2 2001-03-16 17:45:49 11.637444
## 10   2 2001-05-02 04:14:56 10.085473

```

2. For each id, determine if there is more than a one year gap in between observations. Add a new row at the one year mark, with the a1c value set to missing. A two year gap would require two new rows, and so forth.

```

# insert a row after the index indicated
insertRow <- function(data, new_row, index) {
  data_new <- rbind(data[1:index, ],
                    new_row,

```

```

        data[- (1:index), ])
rownames(data_new) <- 1:nrow(data_new)
return(data_new)
}

current_id = 0
first_test = NA
complete_x <- order_x

for (i in (1: nrow(complete_x))) {
  new_id = complete_x$id[i]
  second_test = complete_x$dt[i]

  if (new_id == current_id) {
    n_gap = as.numeric(second_test - first_test) /365.25

    if (n_gap > 1) {

      for (j in (1:n_gap)) {
        complete_x <- insertRow(complete_x, c(NA, NA, NA), i-1+j-1)
        complete_x$id[i-1+j]=current_id
        complete_x$dt[i-1+j]=first_test+as.difftime(365.25 * j, units = "days")
      }

    }

  }

  current_id = new_id
  first_test = second_test
}

head(complete_x, 50)

```

```

##      id      dt      a1c
## 1  1 2001-05-08 16:22:52 7.309995
## 2  1 2001-06-17 22:42:23 8.310721
## 3  1 2001-08-17 16:51:46 6.548845
## 4  1 2001-12-14 14:50:29 5.985275
## 5  1 2002-08-19 13:51:47 6.011547
## 6  1 2003-03-22 03:51:36 7.243858
## 7  1 2003-06-27 01:01:34 5.170870
## 8  2 2001-03-05 22:24:43 9.237660
## 9  2 2001-03-16 17:45:49 11.637444
## 10 2 2001-05-02 04:14:56 10.085473
## 11 2 2001-05-28 12:41:17 11.362266
## 12 2 2001-10-29 11:33:48 8.089224
## 13 2 2001-11-10 11:02:55 9.159491
## 14 2 2002-01-03 05:20:50 7.604405
## 15 2 2002-01-12 04:20:47 8.209176
## 16 2 2003-01-12 10:20:47      NA
## 17 2 2003-06-17 01:43:18 8.743263

```

```
## 18 2 2003-06-26 19:40:59 10.051962
## 19 2 2003-12-05 08:06:49 10.548467
## 20 2 2003-12-28 17:19:13 9.966982
## 21 2 2004-09-19 22:07:42 10.564603
## 22 2 2004-09-20 04:53:12 10.606105
## 23 2 2004-11-27 15:33:28 10.970467
## 24 3 2000-05-01 17:21:57 6.507974
## 25 3 2000-07-04 22:09:43 7.735319
## 26 3 2000-12-24 14:58:33 6.017964
## 27 3 2001-03-29 05:37:39 6.209069
## 28 3 2001-05-26 07:08:17 7.800187
## 29 3 2002-05-26 13:08:17 NA
## 30 3 2002-10-01 08:42:43 6.459650
## 31 3 2003-01-09 11:49:40 8.543998
## 32 3 2004-01-09 17:49:40 NA
## 33 3 2004-01-10 13:37:25 10.047035
## 34 3 2004-03-02 03:03:24 5.551797
## 35 3 2004-06-15 19:14:53 5.541563
## 36 3 2004-07-17 06:47:34 6.055469
## 37 4 2000-05-04 05:40:00 7.892846
## 38 4 2000-06-10 08:40:51 7.871581
## 39 4 2001-03-21 05:55:52 8.264556
## 40 4 2001-08-11 23:41:11 9.045372
## 41 4 2002-02-26 04:44:59 7.255024
## 42 4 2002-09-23 13:23:06 8.667542
## 43 4 2003-09-23 19:23:06 NA
## 44 4 2004-03-12 22:45:37 9.324084
## 45 4 2004-04-24 05:52:05 7.214870
## 46 5 2000-06-03 03:57:21 8.098769
## 47 5 2001-04-07 14:27:32 8.558121
## 48 5 2002-01-13 22:31:45 10.202306
## 49 5 2002-01-15 16:20:01 9.719515
## 50 5 2002-01-21 18:47:11 9.463840
```

3. Create a new column `visit`. For each `id`, add the visit number. This should be 1 to `n` where `n` is the number of observations for an individual. This should include the observations created with missing `alc` values.

```
# id_count <- rep(1, length(unique(complete_x$id)))
#
# for (i in (1:(nrow(complete_x)-1))) {
#
#   current_id = complete_x$id[i]
#   next_id = complete_x$id[i+1]
#
#   if (current_id == next_id) {
#     id_count[current_id] = id_count[current_id] + 1
#   }
# }
#
# for (i in (1:(nrow(complete_x)))) {
#   current_id = complete_x$id[i]
#   complete_x$visit[i] = id_count[current_id]
# }
```

```

#
# head(complete_x, 30)

id_count = 1

for (i in (1:(nrow(complete_x)-1))) {

  current_id = complete_x$id[i]
  next_id = complete_x$id[i+1]

  if (current_id == next_id) {
    complete_x$visit[i] = id_count
    id_count = id_count + 1
  } else {
    complete_x$visit[i] = id_count
    id_count = 1
  }
}

head(complete_x, 30)

```

##	id	dt	a1c	visit
## 1	1	2001-05-08 16:22:52	7.309995	1
## 2	1	2001-06-17 22:42:23	8.310721	2
## 3	1	2001-08-17 16:51:46	6.548845	3
## 4	1	2001-12-14 14:50:29	5.985275	4
## 5	1	2002-08-19 13:51:47	6.011547	5
## 6	1	2003-03-22 03:51:36	7.243858	6
## 7	1	2003-06-27 01:01:34	5.170870	7
## 8	2	2001-03-05 22:24:43	9.237660	1
## 9	2	2001-03-16 17:45:49	11.637444	2
## 10	2	2001-05-02 04:14:56	10.085473	3
## 11	2	2001-05-28 12:41:17	11.362266	4
## 12	2	2001-10-29 11:33:48	8.089224	5
## 13	2	2001-11-10 11:02:55	9.159491	6
## 14	2	2002-01-03 05:20:50	7.604405	7
## 15	2	2002-01-12 04:20:47	8.209176	8
## 16	2	2003-01-12 10:20:47	NA	9
## 17	2	2003-06-17 01:43:18	8.743263	10
## 18	2	2003-06-26 19:40:59	10.051962	11
## 19	2	2003-12-05 08:06:49	10.548467	12
## 20	2	2003-12-28 17:19:13	9.966982	13
## 21	2	2004-09-19 22:07:42	10.564603	14
## 22	2	2004-09-20 04:53:12	10.606105	15
## 23	2	2004-11-27 15:33:28	10.970467	16
## 24	3	2000-05-01 17:21:57	6.507974	1
## 25	3	2000-07-04 22:09:43	7.735319	2
## 26	3	2000-12-24 14:58:33	6.017964	3
## 27	3	2001-03-29 05:37:39	6.209069	4
## 28	3	2001-05-26 07:08:17	7.800187	5
## 29	3	2002-05-26 13:08:17	NA	6
## 30	3	2002-10-01 08:42:43	6.459650	7

4. For each id, replace missing values with the mean a1c value for that individual.

```
id_mean <-complete_x %>%
  group_by(id) %>%
  summarize(mean = mean(a1c, na.rm = TRUE)) %>%
  ungroup()

complete_x_woNA <- complete_x

for (i in which(is.na(complete_x_woNA$a1c))) {
  current_id = complete_x_woNA$id[i]

  complete_x_woNA$a1c[i] = id_mean$mean[current_id]
}

head(complete_x_woNA, 50)
```

##	id	dt	a1c	visit
## 1	1	2001-05-08 16:22:52	7.309995	1
## 2	1	2001-06-17 22:42:23	8.310721	2
## 3	1	2001-08-17 16:51:46	6.548845	3
## 4	1	2001-12-14 14:50:29	5.985275	4
## 5	1	2002-08-19 13:51:47	6.011547	5
## 6	1	2003-03-22 03:51:36	7.243858	6
## 7	1	2003-06-27 01:01:34	5.170870	7
## 8	2	2001-03-05 22:24:43	9.237660	1
## 9	2	2001-03-16 17:45:49	11.637444	2
## 10	2	2001-05-02 04:14:56	10.085473	3
## 11	2	2001-05-28 12:41:17	11.362266	4
## 12	2	2001-10-29 11:33:48	8.089224	5
## 13	2	2001-11-10 11:02:55	9.159491	6
## 14	2	2002-01-03 05:20:50	7.604405	7
## 15	2	2002-01-12 04:20:47	8.209176	8
## 16	2	2003-01-12 10:20:47	9.789132	9
## 17	2	2003-06-17 01:43:18	8.743263	10
## 18	2	2003-06-26 19:40:59	10.051962	11
## 19	2	2003-12-05 08:06:49	10.548467	12
## 20	2	2003-12-28 17:19:13	9.966982	13
## 21	2	2004-09-19 22:07:42	10.564603	14
## 22	2	2004-09-20 04:53:12	10.606105	15
## 23	2	2004-11-27 15:33:28	10.970467	16
## 24	3	2000-05-01 17:21:57	6.507974	1
## 25	3	2000-07-04 22:09:43	7.735319	2
## 26	3	2000-12-24 14:58:33	6.017964	3
## 27	3	2001-03-29 05:37:39	6.209069	4
## 28	3	2001-05-26 07:08:17	7.800187	5
## 29	3	2002-05-26 13:08:17	6.951820	6
## 30	3	2002-10-01 08:42:43	6.459650	7
## 31	3	2003-01-09 11:49:40	8.543998	8
## 32	3	2004-01-09 17:49:40	6.951820	9
## 33	3	2004-01-10 13:37:25	10.047035	10
## 34	3	2004-03-02 03:03:24	5.551797	11

```
## 35 3 2004-06-15 19:14:53 5.541563 12
## 36 3 2004-07-17 06:47:34 6.055469 13
## 37 4 2000-05-04 05:40:00 7.892846 1
## 38 4 2000-06-10 08:40:51 7.871581 2
## 39 4 2001-03-21 05:55:52 8.264556 3
## 40 4 2001-08-11 23:41:11 9.045372 4
## 41 4 2002-02-26 04:44:59 7.255024 5
## 42 4 2002-09-23 13:23:06 8.667542 6
## 43 4 2003-09-23 19:23:06 8.191985 7
## 44 4 2004-03-12 22:45:37 9.324084 8
## 45 4 2004-04-24 05:52:05 7.214870 9
## 46 5 2000-06-03 03:57:21 8.098769 1
## 47 5 2001-04-07 14:27:32 8.558121 2
## 48 5 2002-01-13 22:31:45 10.202306 3
## 49 5 2002-01-15 16:20:01 9.719515 4
## 50 5 2002-01-21 18:47:11 9.463840 5
```

5. Print mean a1c for each id.

```
id_mean$mean
```

```
## [1] 6.654444 9.789132 6.951820 8.191985 9.429694 7.133443 7.879138
## [8] 6.244061 4.420523 6.028370 4.838279 6.691181 8.504632 9.122968
## [15] 6.737092 7.420245 6.546329 6.151311 8.628037 8.923518 5.444430
## [22] 5.763931 6.351112 9.377525 5.058097 8.692078 7.371831 4.243469
## [29] 6.345254 4.135795 8.670622 5.130167 6.528153 8.445030 3.832195
## [36] 9.514603 8.612608 10.160773 8.976697 7.583232 3.804325 6.787170
## [43] 5.654235 5.613283 8.876623 7.485824 4.752133 7.415459 5.562809
## [50] 4.970288
```

6. Print total number of visits for each id.

```
id_count <- rep(1, length(unique(complete_x$id)))
```

```
for (i in (1:(nrow(complete_x)-1))) {
  current_id = complete_x$id[i]
  next_id = complete_x$id[i+1]

  if (current_id == next_id) {
    id_count[current_id] = id_count[current_id] +1
  }
}
```

```
id_count
```

```
## [1] 7 16 13 9 14 11 7 12 15 8 12 12 9 12 10 8 10 14 10 11 13 12 10 12 16
## [26] 11 10 15 3 13 11 9 12 12 11 10 8 14 14 11 14 11 8 12 6 11 9 4 10 7
```

7. Print the observations for id = 15.

```
complete_x_woNA %>%
  filter(id == 15)
```

```
##      id          dt      a1c visit
## 1  15 2000-10-21 01:08:17 7.401322    1
## 2  15 2001-08-08 14:23:08 5.896318    2
## 3  15 2001-08-15 07:03:29 7.457722    3
## 4  15 2002-03-15 21:23:10 5.330917    4
## 5  15 2002-04-14 09:08:25 6.484003    5
## 6  15 2002-10-10 18:27:43 8.139101    6
## 7  15 2003-02-19 12:58:53 6.446557    7
## 8  15 2003-03-02 06:58:10 7.432291    8
## 9  15 2003-06-30 07:20:49 7.113792    9
## 10 15 2004-01-22 20:30:42 5.668897   10
```

## Question 2

### 16 points

Install the `lexicon` package. Load the `sw_fry_1000` vector, which contains 1,000 common words.

```
data('sw_fry_1000', package = 'lexicon')
head(sw_fry_1000)
```

```
## [1] "the" "of" "to" "and" "a" "in"
```

1. Remove all non-alphabetical characters and make all characters lowercase. Save the result as `a`.

```
a <- tolower(gsub("[^a-zA-Z]", "", sw_fry_1000))
```

Use vector `a` for the following questions. (2 points each)

2. How many words contain the string “ar”?

```
sum(grepl("ar", a))
```

```
## [1] 64
```

3. Find a six-letter word that starts with “l” and ends with “r”.

```
a[nchar(a) == 6 & grepl("^l.*r$", a)]
```

```
## [1] "letter"
```

4. Return all words that start with “col” or end with “eck”.

```
a[grepl("^col", a) | grepl("eck$", a)]
```

```
## [1] "color" "cold" "check" "collect" "colony" "column" "neck"
```

5. Find the number of words that contain 4 or more adjacent consonants. Assume “y” is always a consonant.



```
sum(grepl("[bcdfghjklmnpqrstvwxyz]{4}", a, ignore.case = TRUE))
```

```
## [1] 8
```

6. Return all words with a “q” that isn’t followed by a “ui”.

```
a[grepl("q(?!ui)", a, perl = TRUE)]
```

```
## [1] "question" "equate" "square" "equal" "quart" "quotient"
```

7. Find all words that contain a “k” followed by another letter. Run the `table` command on the first character following the first “k” of each word.

```
k_followed <- a[grepl("k[a-zA-Z]", a)]

first_char_after_k <- sapply(k_followed, function(word) {
  # Find the position of k
  pos <- regexpr("k[a-zA-Z]", word)
  if (pos[1] != -1) {
    # Get the character after k
    return(substr(word, pos[1] + 1, pos[1] + 1))
  }
})

table(first_char_after_k)
```

```
## first_char_after_k
## e i n y
## 10 5 2 1
```

8. Remove all vowels. How many character strings are found exactly once?

```
no_vowels <- gsub("[aeiouAEIOU]", "", a)

unique_count <- table(no_vowels)
exactly_once <- sum(unique_count == 1)
exactly_once
```

```
## [1] 581
```

### Question 3

#### 3 points

The first argument to most functions that fit linear models are formulas. The following example defines the response variable `death` and allows the model to incorporate all other variables as terms. `.` is used to mean all columns not otherwise in the formula.

```
url <- "https://github.com/couthcommander/Bios6301/raw/main/datasets/haart.csv"
haart_df <- read.csv(url)[,c('death','weight','hemoglobin','cd4baseline')]
coef(summary(glm(death ~ ., data=haart_df, family=binomial(logit))))
```

```
##              Estimate Std. Error  z value    Pr(>|z|)
## (Intercept)  3.576411744 1.226870535  2.915069 0.0035561039
## weight      -0.046210552 0.022556001 -2.048703 0.0404911395
## hemoglobin  -0.350642786 0.105064078 -3.337418 0.0008456055
## cd4baseline  0.002092582 0.001811959  1.154872 0.2481427160
```

Now imagine running the above several times, but with a different response and data set each time. Here's a function:

```
myfun <- function(dat, response) {
  form <- as.formula(response ~ .)
  coef(summary(glm(form, data=dat, family=binomial(logit))))
}
```

Unfortunately, it doesn't work. `tryCatch` is "catching" the error so that this file can render to PDF.

```
tryCatch(myfun(haart_df, death), error = function(e) e)
```

```
## <simpleError in eval(predvars, data, env): object 'death' not found>
```

What do you think is going on? Consider using `debug` to trace the problem.

In `as.formula`, the function treated `response` as a string but not a variable that holds the string we want to use. So, when calling `myfun(haart_df, death)`, the response variable name "death" stored in 'response' is not used, the function just use 'response' as the response variable name, which causes error since 'response' is not a variable in the dataset. Also, should use `myfun(haart_df, "death")` instead of `myfun(haart_df, death)` as the second argument need to be a string.

## 5 bonus points

Create a working function.

```
myfun <- function(dat, response) {
  regression <- paste0(response, " ~ ", ".")
  form <- as.formula(regression)
  coef(summary(glm(form, data=dat, family=binomial(logit))))
}
```

```
myfun(haart_df, "death") #use string input in the second argument then the function would work
```

```
##              Estimate Std. Error  z value    Pr(>|z|)
## (Intercept)  3.576411744 1.226870535  2.915069 0.0035561039
## weight      -0.046210552 0.022556001 -2.048703 0.0404911395
## hemoglobin  -0.350642786 0.105064078 -3.337418 0.0008456055
## cd4baseline  0.002092582 0.001811959  1.154872 0.2481427160
```